Sequential Learning

Home Assignment

This homework should be uploaded by **Friday, March 15, 2024** as a pdf file on the website

`http://pierre.gaillard.me/teaching/mva2024.php`

The password to upload is `mva2024`. The homework can be done alone or in groups of two students. The code can be done in any langage (`python`, `R`, `matlab`,...) and should not be returned but the results and the figures must be included into the pdf report.

All questions require a proper mathematical justification or derivation (unless otherwise stated), but most questions can be answered concisely in just a few lines. No question should require lengthy or tedious derivations or calculations.

Part 1. Bernoulli Bandits

We consider a stochastic bandit setting in which the arm rewards have Bernoulli distributions. A random variable X is said to have Bernoulli distribution with parameter p, which we denote by $\mathcal{B}(p)$, if it takes value 0 with probability $1-p$ and value 1 with probability p. The set $\{1,\dots,K\}$ is denoted by $[K]$.

Each arm $k \in [K]$ has a reward distribution $\mathcal{B}(p_k)$.
At each round $t = 1,\dots,T$
- The player chooses an arm $a_t \in [K]$,
- The player observes a reward $X_{a_t}(t) \sim \mathcal{B}(p_{k_t})$, independent of all other rewards.

Setting 1: Bernoulli bandit

Notations:
- In this part, the term "regret" refers to the quantity $R_T = \max_{k\in[K]} Tp_k - \sum_{t=1}^T p_{a_t}$.
- N_t^k denotes the number of pulls of arm k immediately after time t, i.e. $N_k(t) = \sum_{s=1}^t \mathbb{I}\{a_s = k\}$.
- $\hat{\mu}_k(t)$ denotes the empirical mean of arm k after time t: $\hat{\mu}_k(t) = \frac{1}{N_k(t)} \sum_{s=1}^t X_{a_s}(s)\mathbb{I}\{a_s = k\}$.

A bit of context: why Bernoulli bandits matter. Many applications have binary outcomes, in which the reward then follows a Bernoulli distribution. A prominent example is online advertising, in which a seller shows advertisements to visitors of a website, and a usual goal is to maximize the probability that the visitor clicks on the ad. In its most basic form, this is exactly the bandit interaction described above: the seller (player) chooses an ad (arm) which is displayed to the visitor, and then the seller observes whether there is a click or not (reward). More elaborate models of that interaction take into account prior information that the seller has about the visitor, turning it into a *contextual* bandit, or get rid of the independence assumption, etc.

1. *Follow the leader.* All experiments in this question will be done for $K = 2$, $p = (0.5, 0.6)$.

(a) Prove that the expected regret of the Follow-the-leader algorithm (`FTL`) verifies $\mathbb{E}R_T \geq \alpha T$, for some $\alpha > 0$. Recall that `FTL` pulls at each time the arm with highest empirical mean: $a_{t+1} \in \arg\max_{k\in[K]} \hat{\mu}_k(t)$.

(b) Implement `FTL`.

(c) For time $T = 100$, plot a histogram of the regret R_T of `FTL` over 1000 repetitions of the experiment. Explain the figure.

(d) Plot the mean regret of `FTL` over 1000 repetitions, as a function of $t \in \{1, \dots, 1000\}$. Is FTL a good algorithm for stochastic bandits?

2. *UCB.* Recall that a random variable is said to be σ^2-sub-Gaussian if for all $\lambda \in \mathbb{R}$, $\mathbb{E}[e^{\lambda(X-\mathbb{E}[X])}] \leq e^{\frac{1}{2}\sigma^2\lambda^2}$. The $\texttt{UCB}(\sigma^2)$ algorithm pulls arm $a_{t+1} = \arg\max_{k\in[K]} \hat{\mu}_k(t) + \sqrt{\frac{2\sigma^2\xi\log(t)}{N_k(t)}}$. It is designed to have low regret on σ^2-sub-Gaussian random variables. For theoretical regret bounds to hold, ξ should be taken slightly larger than 1.[1] **All experiments in this question will be done for $\xi = 1.1$.**

(a) Compute the cumulant generating function, defined for $\lambda \in \mathbb{R}$ by $\phi_X(\lambda) = \log \mathbb{E}[e^{\lambda(X-\mathbb{E}[X])}]$, for a Bernoulli random variable with parameter p.

(b) Prove that if a random variable X (not necessarily Bernoulli) verifies $\phi_X''(\lambda) \leq \sigma^2$ for all $\lambda \in \mathbb{R}$, then the random variable is σ^2-sub-Gaussian. Remark: this is not an equivalence (you are not required to prove this).

(c) Using question 2.b, find σ^2 such that a random variable with distribution $\mathcal{B}(p)$ is σ^2-sub-Gaussian.

(d) Prove that a random variable X supported on $[0,1]$ with mean $p \in [0,1]$ verifies $\phi_X(\lambda) \leq \phi_Y(\lambda)$ for all $\lambda \in \mathbb{R}$, where Y has a $\mathcal{B}(p)$ distribution. Hint: prove that for all $x \in [0,1]$, for all $\lambda \in \mathbb{R}$, $e^{\lambda x} \leq 1 - x + xe^{\lambda}$.

(e) Prove that all random variables supported on $[0,1]$ are $\frac{1}{4}$-sub-Gaussian.

(f) Implement the $\texttt{UCB}(\sigma^2)$ algorithm.

(g) Plot the mean regret of $\texttt{UCB}(1/4)$ as a function of time up to $T = 1000$ for $K = 2$, $p = (0.5, 0.6)$, over 1000 repetitions. Compare with the result of question 1.d.

(h) For $K = 2$, $p = (0.6, 0.5)$, $T = 1000$, plot the mean regret of $\texttt{UCB}(\sigma^2)$ over 1000 repetitions as a function of σ^2, for $\sigma^2 \in \{0, 1/32, 1/16, 1/4, 1\}$. Do it again for $p = (0.85, 0.95)$ and compare the results: does the optimal parameter change? How does it compare to the theoretic parameter?

The results of the question 2.c on Bernoulli distributions can be improved: it is possible to prove that a random variable with distribution $\mathcal{B}(p)$ is σ^2-sub-Gaussian with parameter $\sigma^2(p) = 0$ if $p \in \{0,1\}$, $\sigma^2(p) = 1/4$ if $p = 1/2$ and $\sigma^2(p) = \frac{1}{2}\frac{p-(1-p)}{\log p - \log(1-p)}$ for $p \in (0,1) \setminus \{1/4\}$.

3. On the same figure, plot the variance of $\mathcal{B}(p)$ and the sub-Gaussian constant $\sigma^2(p)$ described above as a function of $p \in [0,1]$.

4. **(optional)** Prove that a σ^2-sub-Gaussian random variable has variance bounded by σ^2.

Adaptation to the variance. The algorithm $\texttt{UCB}(\sigma^2)$ uses only the empirical mean of the arms to choose the next arm, except for a parameter σ^2 which has to be chosen such that all arms are σ^2-sub-Gaussian. In particular, all variance information about the distributions is lost. Intuitively an arm with lower variance should require fewer samples in order to know its mean with enough precision.

[1] We used $\xi = 4$ in class, but better bounds can be obtained as $\xi \to 1$ (with $\xi > 1$).

5. UCB-V. For bounded rewards belonging to $[0,1]$, the algorithm `UCB-V`(ξ, c) (V for variance) computes the empirical variance of the arms, $\hat{v}_k(t) = \frac{1}{N_k(t)} \sum_{s=1}^t \mathbb{I}\{a_s = k\}(X_{a_s}(s) - \hat{\mu}_k)(t)^2$ and pulls the arm $a_{t+1} = \arg\max_{k\in[K]} \hat{\mu}_k(t) + \sqrt{\frac{2\hat{v}_k(t)\xi \log t}{N_k(t)}} + \frac{3bc\xi}{N_t^k}$. Again in theory, ξ should be taken slightly larger than 1 and c larger than a function of ξ, which increases as $\xi \to 1$. **All experiments in this question will be done for $\xi = 1.1$ and $c = 1$.**
   (a) Prove that $N_k(t)\hat{v}_k(t) = \sum_{s=1}^t \mathbb{I}\{a_s = k\}(X_{a_s}(s))^2 - \frac{1}{N_k(t)}(\sum_{s=1}^t \mathbb{I}\{a_s = k\}X_{a_s}(s))^2$.
   (b) Prove that $N_{a_{t+1}}(t+1)\hat{v}_{a_{t+1}}(t+1) = N_{a_{t+1}}(t)\hat{v}_{a_{t+1}}(t) + (X_{a_{t+1}}(t+1) - \hat{\mu}_{a_{t+1}}(t))(X_{a_{t+1}}(t+1) - \hat{\mu}_{a_{t+1}}(t+1))$. What is the practical advantage of that formulation?
   (c) Implement `UCB-V`.
   (d) On the same figure, plot the mean regret of `UCB-V` and `UCB`$(1/4)$ as a function of time up to $T = 1000$ for $K = 2$, $p = (0.5, 0.6)$, over 1000 repetitions.
   (e) Same question for $p = (0.1, 0.2)$ and $p = (0, 0.1)$. Compare to the results of 5.d. When does `UCB-V` improve over `UCB`?

Algorithms for parametric distributions. `UCB` uses only an estimate of the mean, while `UCB-V` uses estimates of the mean and variance. However, Bernoulli distributions have many properties beyond their mean and variance, and these properties are not used by `UCB-V`. We can design algorithms that use fully the knowledge that the distribution of the arms are Bernoulli $\mathcal{B}(\mu)$, with the only unknown being the parameter μ. This is the case of the celebrated `Thompson sampling` algorithm. `Thompson sampling` starts with prior distributions p_k^0 on the parameters μ_k of the bandits instance. At each round t, it draws random samples $\theta_k(t) \sim p_k^{t-1}$ where p_k^{t-1} is the current prior distribution on the parameter μ_k at time $t-1$. `Thompson sampling` then pulls the arm $a_t \in \arg\max_k \theta_k(t)$ at time t. After observing a reward $X_{a_t}(t)$, the prior of the pulled arm is then updated following the Bayes' rule:

$$f_{a_t}^t(x) \propto f_{a_t}^{t-1}(x)\mathbb{P}\left(X_{a_t}(t) \mid \mu_k = x\right),$$

where f_k^t is the density of p_k^t.[2]

6. `Thompson sampling`. Consider in this question **uniform priors** $p_k^0 = \mathcal{U}([0,1])$ and **Bernoulli rewards**.
   (a) **(optional)** Show that the prior distribution p_k^t of the arm k at time t is a Beta distribution with parameters $(S_k(t) + 1, N_k(t) - S_k(t) + 1)$ where $S_k(t) = \sum_{s=1}^t \mathbb{I}\{a_s = k\}X_{a_s}(s)$.
   (b) Implement `Thompson sampling` using the previous question.
   (c) Compare `Thompson sampling` with `UCB` and `UCB-V` in the settings of questions 5.d and 5.e. Comment.

   Remark: an adaptation of `UCB` called `kl-UCB` is also designed precisely to take advantage of the knowledge that distributions belong to a so-called one-parameter exponential family, and that algorithm can use fully the Bernoulli assumption.

Part 2. Rock Paper Scissors

We consider the sequential version of a repeated two-player zero-sum games between a player and an adversary.

Let $L \in [-1,1]^{M\times N}$ be a loss matrix.
At each round $t = 1, \dots, T$
- The player chooses a distribution $p_t \in \Delta_M := \{p \in [0,1]^M, \sum_{i=1}^M p_i = 1\}$
- The adversary chooses a distribution $q_t \in \Delta_N$
- The actions of both players are sampled $i_t \sim p_t$ and $j_t \sim q_t$
- The player incurs the loss $L(i_t, j_t)$ and the adversary the loss $-L(i_t, j_t)$.

Setting 2: Setting of a sequential two-player zero sum game

1. Recall M, N and a loss matrix $L \in [-1,1]^{M\times N}$ that corresponds to the game "Rock paper scissors"[3].

Full information feedback We assume that both players know the matrix L in advance and can compute $L(i,j)$ for any (i,j). Define a function `EWA_update` that takes as input a learning rate $\eta > 0$, a vector $p_t \in \Delta_M$ and a loss vector $\ell_t \in [-1,1]^M$ and return the updated vector $p_{t+1} \in \Delta_M$ defined for all $i \in [M]$ by

$$p_{t+1}(i) = \frac{p_t(i)\exp(-\eta \ell_t(i))}{\sum_{j=1}^M p_t(j)\exp(-\eta \ell_t(j))}\,.$$

2. Consider the game "Rock paper scissors" and assume that the adversary chooses the optimal response to the player: $q_t \in \arg\max_{q\in\Delta_N} \left\{p_t^\top L q\right\}$ and samples $j_t \sim q_t$ for all rounds $t \geq 1$ (note that q_t is likely to be a Dirac distribution).
   (a) What is the loss $\ell_t(i)$ incurred by the player if he chooses action i at time t?
   (b) Simulate an instance of the game for $t = 1, \dots, T = 100$ for $\eta = 1$. Plot the evolution of the weight vectors $p_1, p_2, \dots, p_T$.
   (c) Define $\bar{p}_t = \frac{1}{t}\sum_{s=1}^t p_s$. Plot in log log scale $\|\bar{p}_t - (1/3,1/3,1/3)\|_2$ as a function of $t = 1, \dots, 10\,000$. What do you observe?
   (d) Plot the average loss $\bar{\ell}_t = \frac{1}{t}\sum_{s=1}^t \ell(i_s, j_s)$ as a function of t.
   (e) Repeat one simulation for different values of learning rates $\eta \in \{0.01, 0.05, 0.1, 0.5, 1\}$. What are the best η in practice and in theory?

Bandit feedback Now, we assume that the players do not know the game in advance but only observe the performance $L(i_t, j_t)$ (that we scale here to be in $[0,1]$) of the actions played at time t. They need to learn the game and adapt to the adversary step by step as more feedback is observed.

3. Explain the main differences between EXP3 and EWA and implement `EXP3`.

4. Repeat Questions 2.b) to 2.d) with `EXP3` instead of `EWA`.

[3] This is a common game where two players choose one of 3 options: (Rock, Paper, Scissors). The winner is decided according to the following: Rock crushes scissors, Paper covers Rock, Scissors cuts paper